

# SAPIENS — PRODUCT REQUIREMENTS DOCUMENT

Things Covered:

- Spine
- Safety/SOS
- Identity/KYC/Legal
- Matching & Requests
- Active-Request Comms
- Connection
- Inbox
- Moneta/Meters/Celestial
- Community Moments
- Profile
- UI/UX & Navigation
- Tech Stack & Infrastructure

## HOW TO READ THIS

Phase tags on every decision:

- **[P1]** : Phase 1 / MVP. Ships in the first version.
- **[P2]** : Deliberately deferred; v1 architecture must accommodate it without rework.
- **[LATER]** : Beyond P2. Directional only.
- **[OPEN]** : Not yet decided (tracked in the Open Items Tracker).
- **[GATE]** : Hard pre-launch blocker. Cannot ship publicly without it.

Rationale is recorded next to each decision in the future-you need to know *why*, so decisions are revisited deliberately, never by accident. When later thinking changes an earlier decision, the earlier decision is **updated in place** with a change note (see 0.3 for a live example).

**Phasing is the spine of the plan.** A tight, buildable MVP core first; features added in versioned updates. Every P2/LATER item has a v1 architectural hook.

## PRODUCT DEFINITION

**What Sapiens is:** a hyperlocal mutual-aid app where verified real people help each other in person, for free. No money changes hands. No profiles to browse. No feed to scroll. Trust is built by helping, and recorded in a way that cannot be bought.

**The constitution (non-negotiables):**

1. **No money for help, ever.**
2. **No profile surfing** : information about another person is *earned* step by step.

3. **Everyone is a verified real human** before giving or receiving help.
4. **Trust is the product** : any feature that erodes it is wrong, regardless of engagement upside.
5. **Never endanger the person we're protecting** : safety design beats safety theatre.
6. **We don't farm attention** : no dopamine mechanics, ads, or engagement traps.

## Vocabulary :

| Term                       | Meaning                                                                                                                                                                                     |
|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Moneta</b>              | Token for each first-ever (unique) help to a person. Flat: 1 per unique help. Non-transferable. No published cash value in v1; redeemable for goodies later.                                |
| <b>Goodness Meter</b>      | 0–100 speedometer from unique helps, responsive curve, lifetime climb. Owner-facing. The "width of your circle."                                                                            |
| <b>Celestial Journey</b>   | Growth badge (= the Spiritual Meter): new moon → crescent (10) → half (50) → full (100) → sunrise (500) → golden sun (1000) → joins the galaxy. Driven by unique help. Public trust signal. |
| <b>Connection</b>          | Mutual, opt-in relationship formed only after a real completed help.                                                                                                                        |
| <b>Seeker / Helper</b>     | The person who raised the request / the person who responds.                                                                                                                                |
| <b>Request timing</b>      | Every request is Now (live, expiring) or Scheduled (future, patient). New in Bucket 3.                                                                                                      |
| <b>Active-request chat</b> | The temporary coordination chat between matched parties. Scaffolding, not a room, it dissolves at completion. New in Bucket 4.                                                              |
| <b>Directed request</b>    | A help request aimed at a specific connection rather than open dispatch. New in Bucket 5.                                                                                                   |
| <b>Inbox</b>               | The permanent conversation space, available only between connections. Distinct from the active-request chat. New in Bucket 6.                                                               |
| <b>Unique help</b>         | A first-ever help to a specific person. The basis of ALL rewards/KPIs. Repeat helps the same person earn nothing. New in Bucket 7.                                                          |
| <b>Steadfast count</b>     | The separate, honoured tally of <i>total</i> helps including repeats — recognises depth (helping the same person often). Does not drive Moneta/Meters. New in Bucket 7.                     |
| <b>Community moments</b>   | The finite, area-scoped surface of real completed-help moments (selfies, milestones). Has a bottom; no infinite scroll. Sapiens' answer to "a feed," without the trap. New in Bucket 8.     |

| Term                             | Meaning                                                                                                                                                                             |
|----------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| KYC selfie vs. display photo     | Two separate photos: the locked, private liveness selfie and the editable, public display photo (recognition), every display photo face-matched to the KYC selfie. New in Bucket 9. |
| Premium Choice                   | Paid upgrade letting a seeker pick their helper from a list. [P2]                                                                                                                   |
| Founding Sapiens / City Saviours | First 1000 volunteers (website) / future paid helper roles [LATER].                                                                                                                 |

## BUCKET 0 — THE CORE SPINE

### 0.1 Matching model - DISPATCH [P1]

Dispatch with a confirmation gate, **not** a browsable feed. Seeker raises a request → system pings eligible nearby helpers opted into that category → helpers "raise a hand" → one is confirmed → only then is location/chat/identity revealed.

**Rejected:** a browse-and-claim **feed** (it's the scroll-trap Sapiens exists to escape; strips seeker control - a predator could claim a vulnerable person's request; fails when nobody's looking).

**Seeker-browses-and-picks-first** (exposes profiles up front; slow; burdens a distressed/low-tech seeker). **Hook:** the confirm step means the "choose a helper" UI exists in v1; Premium Choice (0.9) just moves selection one step earlier.

### 0.2 Request tiers [P1]

Two independent tag dimensions per request:

**Category tags** route to the right helpers (a helper is only pinged for categories they opted into). Seed parents: **Travel · Food · Knowledge · Time · Hands · Material · Activity · Safety** (full taxonomy in 3.11).

**Urgency tiers** set ordering, reach, and accept model: **Casual · Everyday · Urgent · SOS/Emergency**.

### 0.3 Accept model — REVISED in Bucket 3 [P1]

 **CHANGE NOTE:** v0.1 set "seeker-picks for Urgent+." Bucket 3 revised this, seeker-picking became the paid **Premium Choice** (0.9), so the free tier could not use it for Urgent. The current model below supersedes the original.

**Current free-tier matching (see 3.6 for full logic):**

**Casual / Everyday** → first-to-accept.

**Urgent** → system **auto-assigns the best single helper** (nearest + reliable), asked one at a time.

**Every non-emergency match** → the seeker sees who accepted (name, photo, badge) and can **decline that specific helper** before location/chat unlocks — a lightweight veto, not a browse-and-pick.

**Seeker-picks-from-a-list** = Premium Choice [P2].

**Rationale:** preserves speed for trivial help, gives Urgent a deliberate (not free-for-all) match, and keeps a last line of seeker control (the veto) consistent with identity-only safety, *you can't browse, but you can refuse*.

#### 0.4 Staged disclosure – the information ladder [P1]

Core principle: **information is earned, step by step**.

| Stage                                               | What's revealed                                                                                                                     |
|-----------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------|
| <b>Request dispatched</b>                           | Helpers see help type + <b>approximate area</b> + urgency. Not who, not exactly where.                                              |
| <b>Helper raises hand</b>                           | Seeker sees helper's <b>first name, verified photo, distance, Celestial badge</b> (1.6). Enough to judge — not a browsable profile. |
| <b>Seeker confirms</b> (or veto→next)               | Mutual reveal: first name + verified photo. <b>Precise location shared</b> . In-app chat opens. Live tracking begins.               |
| <b>At meetup</b>                                    | <b>Meetup code</b> confirms the right person.                                                                                       |
| <b>After help</b>                                   | Rate → optional Connection → <b>full profile + lasting chat only if both opt in</b> .                                               |
| <b>SOS overrides this entire ladder</b> (Bucket 1). |                                                                                                                                     |

#### 0.5 Location policy [P1]

**Approximate during matching** → **precise only after the seeker confirms**. A seeker's home is never revealed before *they* choose the person.

**Live tracking** of the approaching helper once confirmed.

**Public meetpoint default** for casual help.

**Share-my-meetup [P1]:** one tap to send a live-location link to a trusted contact for any ordinary meeting. Reuses SOS trusted-contact plumbing (1.2) nearly free, strongest everyday-safety feature per unit of effort.

#### 0.6 Communication during a request [P1]

**In-app chat only in v1** — text + limited photos/video for context.

**No real phone numbers are ever exchanged during a request**, this resolves the entire masked-number question set (the way to protect numbers is to not exchange them).

**Masked calling** → [P2] (per-call cost; telephony vendor).

Real contact details become possible only if two people **connect** and choose to share. Chat lifecycle/media/moderation/retention → **Bucket 4**.

## 0.7 Arrival & no-shows [P1]

**Live status:** Accepted → On the way → Arrived.

**Auto-timeout:** no "On my way" within X min → acceptance lapses, request **silently re-broadcasts** (kills ghost-acceptances that block others).

**One-tap re-broadcast** when a seeker is let down (they're disappointed *and* still need help — no starting over).

**The app must notice silence:** an unfilled/abandoned request never shows a dead spinner — it offers a widened radius / safe public wait spot / SOS / nudges a trusted contact.

**No-show policy = grace-then-consequence:** first no-show is a shrug; a *pattern* reduces dispatch reach, framed as **rebuilding trust**, not punishment. (Purely punitive fights the brand; purely gentle makes ghosting free.)

## 0.8 Completion, Moneta & rating [P1]

**Completion model:** Helper marks done → seeker confirms → OR auto-confirms after a window → Moneta releases → seeker keeps a dispute window (claw back Moneta + flag). **Why:** the **meetup code already proves they met**, so completion can't begin without a real meeting → the honest helper isn't punished by a silent seeker (the low-tech/elderly/distressed seeker this app serves may never tap "confirm"), and the helper can't farm alone. (Seeker-only rejected: dies on the unresponsive seeker. Helper-only rejected: fraud surface.) **Moneta:** flat **1 per completed help, always, including SOS**; releases only after confirmation; **decoupled from rating** (releases whether or not anyone rates). Keeping Moneta "dumb and honest" makes it ungameable by chasing high-value tasks; nuance lives in the Goodness Meter. **Rating: mutual and double-blind** (stars + optional text/voice; neither sees the other's until both submit) — double-blind kills retaliatory rating.

## 0.9 Premium Choice – "pick your helper" [P2 / monetization]

The rejected seeker-selects model returns as a **paid upgrade**: free users get safe dispatch; paying users choose their helper up front. Philosophically safe — charging for *convenience/control*, never for help itself. Addresses funding (A6) without violating the constitution. Hook: reuses v1's confirm-screen UI, selection moved earlier.

## 0.10 The spine, in one line

Request (category + urgency + Now/Scheduled) → Dispatch to eligible nearby helpers → Hands raised → Match (first-accept / auto-assign for Urgent, with seeker veto) → Confirmed (reveal + precise location + chat + live tracking) → Meet (public point, meetup code) → Help → Complete (helper done, seeker confirms or auto-confirm) → Moneta (flat 1) + double-blind rating → Optional mutual Connection. **SOS branches off at "Request" and skips every gate.**

## BUCKET 1 – SAFETY & SOS

### 1.1 Safety model – identity-only (deterrence) [P1]

v1 = **accountability through verified identity**. Everyone is a real, government-ID-verified, traceable human; misbehaviour → identified and banned. **Acknowledged limitation:** this is **deterrence, not prevention** — KYC proves identity, not safety. A conscious, documented MVP tradeoff. **Consequence:** verification strength is now the *only* thing between a predator and a vulnerable seeker → raises the bar for Bucket 2 (liveness is non-negotiable). **Deferred:** behavioural trust gates (new-user limits, capped radius, enforced public meetpoints until a track record exists) → **[LATER]**; formal criminal background checks → **[LATER]** (no clean consumer criminal-record API in India).

### 1.2 SOS – phased safety system

Original vision ("press → everyone within 5km + nearest police station notified with live location") **redesigned**, because as specified it could endanger the person it protects: (1) broadcasting a victim's precise live location to nearby strangers can expose her further; (2) "notify the police station" isn't buildable — police expose no consumer API; India's real channel is **112 (ERSS)**; claiming police are notified when they aren't is dangerous; (3) it's the widest-reach, no-gate surface — top abuse/liability vector, and sending civilians toward violence endangers responders too.

**Layer 1 — Trusted contacts [P1] (backbone):** user names **1–3 emergency contacts** at onboarding; SOS instantly alerts them + **a live-location link by SMS** (no app needed). People who know you and *will* act — what every credible safety app is built on. **Layer 2 — One-tap 112 [P1]:** SOS offers one-tap dial to 112 with location ready to relay. Honest, legal, real. Formal police integration → **[LATER]** (needs government partnership). **Layer 3 — Community responders [P2] (signature, deferred):** the "3B" constrained shape — only **opt-in, verified, extra-vetted** responders alerted, **tight radius, approximate location first**, precise only on accept. Deferred because it's what can't be gotten wrong on day one, it *needs density* to be a real rescue, and it carries the heaviest liability. Hook: shares the v1 SOS trigger.

**Permanent SOS guardrails (all phases) [P1]:** (1) never broadcast a victim's precise live location to strangers indiscriminately; (2) never turn users into vigilantes — response means *"be a visible safe presence and call the authorities,"* never intervene; (3) SOS stays out of the points economy — no bounty on an emergency; (4) never overstate — users are told plainly SOS alerts *their contacts* and connects to 112, no implied police dispatch.

### 1.3 SOS & Moneta – flat, plus honour [P1]

SOS response earns the **same flat 1 Moneta**; exceptional acts get a **rare, non-fungible, deed-based honour** (recognition, not currency). **Why a 5× SOS reward was rejected** (proposed, then reversed): it puts a bounty on emergencies (motive to raise a false one — a cried-wolf SOS is worse than none); it attracts the wrong responder (points-sprinter vs. calm helper); it compounds the fake-SOS-as-lure attack; and it reopens scaled rewards generally, sliding Moneta back into point-farming.

### 1.4 Everyday meeting safety [P1]

Inherited: KYC before helping · verified-photo reveal at match · meetup code · approximate→precise location · public-meetpoint default · double-blind ratings · dispute/flag window · share-my-meetup.

**Added: Prefer-women helpers [P1] for women seekers** (values + safety, appropriate to context). **In-app reporting + real human response loop [P1]** — visible **Report** on every interaction, backed by human moderation that reviews and acts (not just a button → operational design [OPEN A2]). **Deferred:** forced public-meetpoints for higher-risk requests → [LATER].

### 1.5 Anti-abuse – parked deliberately

**Collusion** (two friends farming Moneta by "helping" each other) — completion can't stop it (they genuinely met). **Approach [P2]:** track **repetition between the same pair**, diminishing returns, pattern detection once Moneta has value. Not built in v1 by decision — *design for the 95% who are good, catch the rest quietly*. Other vectors for the anti-abuse pass: fake requests, urgency inflation, retaliatory rating (mitigated by double-blind), leaderboard gaming, and **the fake-request-as-lure attack** (most serious).

### 1.6 The Celestial badge as trust signal [P1]

The moon→sun badge is shown on a helper's profile and **is the trust indicator the seeker sees at match** (0.4). Elegant because: it makes the engagement metaphor do real **safety** work; it **cannot be bought** (earned only by helping); it's a **signal, not a gate** — never blocks anyone, fully consistent with identity-only. [OPEN A3] **new-helper problem:** a new good person shows as "new moon" and may be skipped — cold-start-of-trust. (*Partially addressed in 3.4: badge does not affect dispatch order, so new helpers still get pinged fairly; the "seeker skips new moon at confirm" half remains open.*)

---

## BUCKET 2 — IDENTITY, KYC & LEGAL

### 2.1 Onboarding & verification gate [P1]

Sign up with **email + phone, OTP-verified**. A new user may **look around** before verifying.

**Government-ID KYC required before giving OR receiving help** — the hard gate. Framing: "Verification isn't a hurdle. It's the foundation."

## 2.2 Verification method [P1]

**Decision: a third-party KYC provider — Aadhaar-primary (via DigiLocker/offline QR under the hood), DL fallback, + liveness selfie. Store only minimal verified fields. Two non-negotiables (forced by identity-only, 1.1):** (1) **liveness selfie matched to the ID photo** — else a borrowed/bought ID makes "traceable" point at the wrong human and the safety model has a hole (bonus: same photo shown at match); (2) **never store a raw Aadhaar number** — store only **verified:yes · name · photo · over-18 flag · verification token**. Resolves the privacy paradox: enough identity for safety, no honeypot, far less legal exposure. **Why a provider (not direct DigiLocker) for MVP** (reverses initial instinct): direct DigiLocker/offline-Aadhaar needs an approved entity + formal onboarding and leaves you building liveness yourself; a provider delivers Aadhaar + DL + liveness behind one integration and returns the minimal payload; faster, lower legal surface, standard for Indian startups. **Direct DigiLocker → [LATER]** cost optimization. ⚠️ **SEQUENCING DEPENDENCY:** most KYC providers require a **registered business entity**. KYC is the whole safety model → app can't ship without KYC → KYC likely can't ship without an entity. **The legal-entity decision may not be fully deferrable.** Everything else builds in parallel meanwhile. **[OPEN A4]** choose vendor (needs fresh research).

## 2.3 Child mode / under-18s [P1 policy / [GATE] legally]

Under-18s → **separate child mode: group-only connections, no one-to-one meetings with unknown adults**. Age from KYC's over-18 flag. **POCSO is strict** — age-gating/enforcement reviewed by counsel before launch (2.4 G4).

## 2.4 🔒 PRE-LAUNCH LEGAL GATE [GATE]

Decided after MVP + initial market testing, but **blocks public launch**. Nothing ships to the public until each is resolved with professional counsel.

| #  | Item                                          | Why it blocks launch                                                                     |
|----|-----------------------------------------------|------------------------------------------------------------------------------------------|
| G1 | Legal entity (company vs Section 8 nonprofit) | Affects KYC access (2.2), liability, funding routes (CSR needs specific structures)      |
| G2 | Liability posture + Terms of Service          | Sharper now safety = deterrence; when prevention fails, "was Sapiens negligent?" is live |
| G3 | DPDP Act (2023) data duties                   | Holds location, identity, chat, verification data on Indian citizens                     |
| G4 | POCSO + minor protection                      | Strictest Indian liability; child-mode age-gating must be legally sound                  |



| #                                                                                                                                                                                    | Item                                                                 | Why it blocks launch                                                                                                                                            |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| G5                                                                                                                                                                                   | IT Rules (platform hosting messaging)                                | Grievance officer, takedown timelines, traceability                                                                                                             |
| G6                                                                                                                                                                                   | Insurance                                                            | When a physical-world interaction goes wrong                                                                                                                    |
| G7                                                                                                                                                                                   | <b>Moneta redemption regulation</b> (once real-world value is added) | Declared/implicit monetary value → RBI prepaid-instrument / voucher / tax exposure.<br><b>Redemption must not launch before the anti-collusion layer (7.4).</b> |
| <b>Owner action:</b> engage a lawyer with Indian tech/platform experience — not self-servable. (Claude can research the landscape + KYC vendors to prepare; cannot replace counsel.) |                                                                      |                                                                                                                                                                 |

## BUCKET 3 — MATCHING ENGINE & REQUEST MODULE

The spine says *dispatch, then confirm*. This bucket sets the mechanics that make it work — and notifications are the real engine (the thing most likely to quietly kill the app if mishandled).

### 3.1 Availability — IMPLICIT [P1]

A helper is reachable **by default unless mid-help, snoozed, or in quiet hours**. No "go online" toggle, no mandatory shift hours. **Rejected:** explicit toggle (people forget it; a *stale* toggle is worse than none) and scheduled-hours-only (rigid; helping is spontaneous). **Rationale:** the ride-hailing analogy misleads — a driver *works*; a neighbour who might share a tiffin is not on shift, and "clock in" turns kindness into a job. At launch density, an availability toggle is a silent killer — only those who remembered to toggle on are eligible. Implicit **maximizes reach exactly when you have the fewest people**. Give **snooze + quiet hours** instead (the opposite default of Uber's).

### 3.2 Radius — ADAPTIVE with a helper cap [P1]

Starts tight, **widens automatically if nobody responds**, up to a **helper-set maximum** (their exposure cap, e.g. "never ping me beyond 5km"). Urgency modulates: SOS opens widest, casual stays tight. Solves cold-start gracefully (a request in a thin area reaches further rather than dying at 3km).

### 3.3 Wave strategy — WAVES, urgency compresses timing [P1]

**Category opt-in is the primary filter — a helper is only ever a candidate for categories they signed up for.** Among candidates, dispatch goes in **waves**: nearest/best few first; if nobody bites in N minutes, widen to the next wave. **Urgency compresses the wave timing** — SOS  $\approx$  instant/all; casual = patient waves. **Rationale:** waves protect the scarcest resource — **helpers' willingness to keep notifications on.** Ping-all dogpiles (twelve pinged for one badminton game, eleven feel useless) and drives the fatigue that makes people mute the app — the app dies quietly with no bad review. **[OPEN A12]** exact wave sizes + intervals per category/urgency need real-world tuning.

### 3.4 Ordering within a wave — nearest-first + rotation; NO badge [P1]

Base order = **nearest-first**, corrected by **load-rotation for fairness** (so the closest active helper isn't hammered into muting while quieter neighbours two streets over are never asked and churn). **The Celestial badge / reliability does NOT affect dispatch order** — dispatch stays fair; the badge informs the *seeker's* choice at the confirm step, never the *system's* ping order. **Rationale:** ranking by badge creates rich-get-richer (high-badge helpers asked first, help more, climb higher; new moons never asked) — directly re-triggering the cold-start-of-trust problem (A3). Fair dispatch is what keeps a *community* alive rather than a handful of saints doing everything.

### 3.5 Request timing — NOW vs SCHEDULED (third dimension) [P1]

Alongside category and urgency, every request is **Now** or **Scheduled** — a *structural* distinction, not fast-vs-slow:

- **Now:** live, expiring, seeker waiting; quick dispatch, response windows.
- **Scheduled:** has a *when*; can sit for hours/days gathering hands; leisurely dispatch + a reminder before the appointed time. **Both ship in v1** — Scheduled is what makes the social/companionship categories (badminton, English, hiking) work, and those are the **safest, richest cold-start fuel**.

### 3.6 The MVP matching model (consolidated) [P1]

- **Casual / Everyday** → first-to-accept.
- **Urgent** → system auto-assigns the best single helper (nearest + reliable), asked one at a time.
- **Every non-emergency match** → seeker sees who accepted (name, photo, badge) and can **decline that specific helper** before location/chat unlocks (veto, not browse).
- **Scheduled** → seeker declares **1-to-1 or group** when raising the request.
- **Premium Choice** (seeker picks from a list) = **[P2]**. (*This is the resolution of the 0.3 collision — see the change note there.*)

### 3.7 Request lifetime [P1]

For **Now** requests: **system sets a sensible default window per urgency** (indicative: casual ~2h, everyday ~1h, urgent ~30m), **seeker can extend**. One less decision when they just need help. Scheduled requests live until their scheduled time (+ grace), then expire. **[OPEN A13]** exact default windows to tune.

### 3.8 Notification budget — the survival mechanic [P1]

- **Soft daily cap** on how many help-request pings one person receives (an active helper in a dense area getting 20/day will mute within a week — and a mute is permanent churn). **SOS is always exempt** from the cap.
- **Batching** of low-urgency pings into a **gentle digest**, not a barrage. **Rationale:** helper attention is the scarce resource; this boring mechanic is *the* thing keeping the network alive. [OPEN A14] exact cap number + digest cadence to tune.

### 3.9 Seeker notifications — two distinct moments [P1]

- **First-to-accept (Casual/Everyday)** → "*X is coming*" (matched, done).
- **Urgent auto-assign** → "*X has accepted — confirm or decline*" (the veto moment; action needed).
- **Dead request** → honest status + "*notify me when someone can help*" (3.10). Each help type has its own response rhythm (some instantaneous, some after time); the seeker is always pinged the moment their request's state changes.

### 3.10 The dead request — fail honestly [P1]

Adaptive widen → prompt seeker to broaden category/timing → then, if still nothing, tell them plainly "*no one's available right now*" with a graceful path: **"notify me when someone can help,"** retry later, or widen the search. **Never a dead spinner, never a pretense.**

### 3.11 THE REQUEST MODULE — taxonomy, schema & architecture [P1]

Architecture principle: categories are DATA, not code

Categories **will** change with demand. The request module reads the taxonomy from config/database, editable **without an app update**. Every category has an **enabled** flag and a **version**, and carries the attributes that drive engine *behaviour* — the category isn't just a label, it tells the engine how to treat the request.

Category schema (per node)

```
{  
  
  id: string,  
  
  parentId: string | null,      // null = top-level parent  
  
  label: string,              // "Ride somewhere"  
  
  slug: string,  
  
  enabled: boolean,          // hide without deleting  
  
  version: number,
```

```

icon: string,           // doodle key (brand-consistent)

description: string,

// --- behaviour-driving attributes ---

defaultTiming: 'now' | 'scheduled' | 'either',

typicalUrgency: 'casual' | 'everyday' | 'urgent', // default tier

interactionType: 'one_to_one' | 'group' | 'either',

impliesPhysicalMeet: boolean, // nearly always true; false reserved for future remote help

safetySensitivity: 'low' | 'medium' | 'high', // drives meetpoint enforcement, women-helper relevance,
extra care

allowsMedia: boolean, // context photos/video useful in chat

radiusHint: 'tight' | 'normal' | 'wide' | null // optional starting radius nudge
}

```

Two-level structure is **stored** now; the UI shows only the **8 parents at launch** (subtypes turn on when volume justifies — decision from this bucket). Categories are **surfaced adaptively** — an area only shows categories it actually has supply for, so an empty "Tennis partners" never looks dead.

Seed taxonomy — 8 parents + subtypes (v1)

Parent-level defaults shown in brackets [timing · typical urgency · interaction · safety].

**1. Travel & Mobility** — [either · everyday · one\_to\_one · high] ride somewhere · accompany on a journey · shared trip (hitchhike-style) · help navigating public transport · pickup (station/airport).

**2. Food & Essentials** — [now · everyday · one\_to\_one · medium] a meal shared · groceries dropped · medicine picked up · water/supplies in a pinch · cooked food for someone unwell.

**3. Knowledge & Skills** — [scheduled · casual · one\_to\_one · medium] language practice · tutoring/homework help · resume/job guidance · tech help (phone setup, filling an online form — the elderly use case) · a professional skill shared.

**4. Time & Presence** — [either · everyday · one\_to\_one · high] companionship for the elderly/lonely · a companion to an appointment · sitting with someone · helping a newcomer find their feet.

**5. Hands & Effort** — [either · everyday · either · medium] moving/lifting · small repairs · assembling something · an extra pair of hands.

**6. Material & Sharing** — [scheduled · casual · one\_to\_one · low] donate an item someone needs · lend a tool · pass on furniture/books/clothes (OLIO-style give-and-take).

**7. Activity & Community** — [scheduled · casual · group · medium] sports partner/group (badminton, tennis, running) · walking companion · hobby groups · someone to do a thing *with*. (*Mostly Scheduled + often group — safest, richest cold-start fuel.*)

**8. Safety & Urgent Help** — [now · urgent/sos · one\_to\_one · high] **SOS** (its own protected system, Bucket 1) · urgent non-emergencies (stranded, locked out, need a hand fast).

*The taxonomy is a starting point — expanded/pruned by demand and user suggestion via the editable config, never a hardcoded app change.*

---

## BUCKET 4 — ACTIVE-REQUEST COMMUNICATION

The temporary chat between matched parties, before any Connection exists. Covers original questions Q15, Q26–30, Q50–59.

### 4.1 The frame — chat is scaffolding, not a room [P1]

The active-request chat has exactly one job: **get two strangers to a successful meeting and a completed help**. Then it goes away. **Rationale:** a rich, persistent chat *before* Connection would (a) quietly violate "connection must be earned" — people would just keep talking and never need to connect, (b) turn Sapiens into a messaging app, which it is not, and (c) accumulate sensitive conversation data the platform would rather not hold. Everything in this bucket follows from that frame: **minimal, functional, ephemeral**.

### 4.2 Lifecycle — opens at confirm, closes at confirmed completion [P1]

- **Opens:** the moment the seeker confirms a helper (per the disclosure ladder, 0.4). Never before — helpers being pinged cannot message the seeker.
- **Closes:** at **confirmed completion** — i.e. when the seeker confirms, or the auto-confirm fires (0.8). **Not** the instant the helper taps "done."

**Why "confirmed," not "marked done":** it leaves the chat live through the confirm window, which acts as a **natural grace period** ("I'm pulling up now," "which gate?", "I left my umbrella") without building an explicit one — and critically, the chat is still open if the seeker needs to say "*wait, this isn't actually done.*" After that point the room dissolves; any further contact requires a **Connection** (Bucket 5). On-brand: continued access to another person is *earned*, never a default.

- **Group requests:** the chat closes when the activity is completed/expired, same principle.

#### 4.3 Data fate — hidden at close, briefly retained for safety only [P1]

On close: messages are **immediately hidden from both users**, then **retained for a short window, access-controlled, for Trust & Safety only**, then permanently deleted. Indicative window: 30–90 days [OPEN A15].

**Why not instant hard-delete** (the purer privacy stance, considered and rejected): if someone is threatened, harassed, or scammed in that chat and reports it, **instant deletion means there is no evidence and the report is worthless**. Retention is what makes the reporting loop (1.4) actually enforceable. **This dovetails with the dispute window (0.8)**: a seeker who reports immediately after completion still has evidence available to a human reviewer. Users get privacy (they can never see or retrieve it); safety gets a workable evidence window; nothing is kept permanently. Ties to **A10** (data retention) and **A2** (moderation ops). Retention rules must be stated plainly in the privacy policy → [GATE G3, DPDP].

#### 4.4 Capabilities — what the chat can carry [P1]

**Allowed**: text · photos · **voice notes** · emoji. **Not in v1**: video, GIFs, arbitrary file uploads.

##### Rationale:

- **Photos** earn their place — the item being donated, the broken thing, the meeting spot.
- **Voice notes are quietly the most important item here**. For the elderly or low-tech seeker who struggles to type — precisely the person this app exists to serve — *talking is far easier than typing*. This is an **accessibility feature, not a frill** (ties to A8).
- **Video/GIFs excluded**: heavier, larger abuse surface, and GIFs specifically pull the tool toward "social/fun," which fights the scaffolding frame (4.1).
- File size limits and media handling: standard sane caps [OPEN A15].

#### 4.5 Moderation — reactive, human-backed [P1]

- **Model: reactive**. A visible **Report** on the chat (and every interaction, per 1.4) sends the conversation to **human review**. **No automated scanning of private conversations** in v1.
- Backed by a **published code of conduct** defining what is banned: harassment, threats, hate speech, sexual content or solicitation, scams, and **any attempt to introduce money into a help interaction** (constitution violation).
- **Automated/proactive filtering** → [LATER].

**Rationale**: proactively scanning private 1-to-1 conversations is heavy to build, privacy-invasive, error-prone, and cuts against the platform's ethos. The report-button-plus-human-loop already committed to *is* the enforcement mechanism — its strength depends entirely on the moderation operations still to be designed (**A2**), and on retention (4.3) making evidence available.

#### 4.6 Contact sharing — allowed, never facilitated [P1]

- **The system never shares or auto-exchanges real phone numbers.** There is no "share my number" button. Numbers become shareable only via a **Connection** (Bucket 5), and even then only if both choose.
- **But the chat does not block or police two adults who choose to exchange details themselves.**

**Rationale:** the principle is about **what the system does**, not about policing what two people say to each other. Technically blocking contact-sharing is paternalistic, brittle (trivially circumvented), and treats users as suspects — the opposite of a trust-first product. **Masked calling remains [P2]** (0.6) — per-call cost + telephony vendor dependency.

#### 4.7 Group coordination chat [P1]

For **Scheduled group** requests (badminton with four, a hiking group), there is no single matched helper — a group forms. This is the original "a dynamic group chat is created" vision, in its correct place.

- The coordination chat is **multi-party: the seeker + every confirmed participant**.
- **People enter as they join and drop out if they leave.**
- The seeker sets a **participant cap** when raising the request ("I need up to 4").
- **Same lifecycle (4.2), same capabilities (4.4), same moderation (4.5), same data fate (4.3)** as the 1-to-1 chat.

#### 4.8 Hand-off at completion [P1]

- If the two parties **Connect** (mutual opt-in, Bucket 5): **the conversation carries over into the permanent inbox**. Continuity feels natural — the relationship began here.
- If they **do not connect**: the chat closes per 4.2 and follows the data policy in 4.3.
- Full inbox behaviour (encryption, blocking, disappearing messages, storage) → **Bucket 6**.

---

## BUCKET 5 — THE CONNECTION PROTOCOL

How a completed help becomes a permanent, opt-in relationship. Covers original questions Q31–Q49, Q107–Q108.

#### 5.1 The frame — the only relationship in Sapiens [P1]

A **Connection** is the single relationship type in the app, and it must be **earned by a real act of help**. There is no way to connect with someone you have not helped or been helped by — no browsing, no friend requests, no discovery. This is the philosophical payoff of the whole product: *unlike every other social app, a relationship here rests on something real that already happened*.

## 5.2 The offer moment [P1]

- The Connect offer appears **immediately on the completion screen**, to both parties, right after rating.
- **Rationale:** warmth is highest at that exact moment — an hour later the feeling has cooled.
- **Rating-gated suppression:** if either party rated the other poorly, the offer is **quietly not shown to either side**. No awkwardness, and — importantly — it leaks no information about what the other person rated. (*Carried from the Q1 proposal; flag if you want this changed.*)

## 5.3 Decline, expiry & resend [P1]

- **Decline is silent.** If one accepts and the other declines, the acceptor is **never told they were declined** — the offer simply lapses. *Nobody is ever informed of rejection.*
- **Offers expire after ~7 days**, then quietly disappear.
- **Resend rule:** a new offer becomes available only if **the two help each other again**. One act = one offer.
- **Rationale:** ties connection permanently to helping, and structurally prevents pestering — someone cannot repeatedly ask to connect with a person who isn't interested.

## 5.4 What unlocks on Connection [P1]

Mutual acceptance unlocks, for both parties:

1. **Full profile visibility** — the end of the disclosure ladder (0.4).
2. **Permanent inbox chat**, with the active-request conversation **carried over** (4.8).
3. **Milestone notifications** about each other (5.8).
4. **Directed requests** — the ability to ask this person for help specifically (5.5).

## 5.5 Directed requests — asking a connection [P1]

(Original Q108.)

- **Routing: directed first, then fallback.** The request goes to that connection alone; if they don't respond within a window, it **falls back to normal open dispatch**. **Rationale:** without fallback, a seeker who asked one friend sits waiting while the whole neighbourhood was available and willing.
- **Disclosure ladder collapses.** No veto step, no meetup code, immediate location/chat — *the ladder exists to protect strangers, and these two are not strangers.*
- **Moneta: yes.** Directed help is real help and earns the flat 1 Moneta like any other.
- **⚠️ Anti-abuse note:** directed requests are the **easiest collusion path in the app** (two friends, repeated "helps," no stranger-facing friction). This is the pair-repetition case parked at [P2] (1.5). **When diminishing-returns logic is built, directed requests between the same pair are the first place it must apply.**

## 5.6 Connections in normal dispatch — the connections wave [P1]

When someone raises a **general** (non-directed) request, their **connections are pinged as a small "connections wave" before wave 1** of nearest-first dispatch (3.3/3.4).



**Rationale:** warmer, more likely to be accepted, and safer than a stranger — cheap to build, large emotional payoff. *(Original Q107.)*

#### 5.7 Privacy of the connection graph [P1]

- **Connection count and list are hidden from strangers and non-connections.**
- **Visible to:** the owner, and their own connections.
- **Rationale:** a public connection count turns relationships into a **follower score** — the exact vanity metric Sapiens exists to reject. The Celestial badge already displays what should matter: *helping*, not popularity. *(Original Q47.)*

#### 5.8 Milestone notifications [P1]

- When a connection reaches a milestone (1, 3, 10, 50, 100, 500, 1000 helps), their connections are notified.
- **Mutable per person** — a user can turn off milestone notifications for any individual connection.
- **Separate from the help-request notification budget** (3.8/A14) — celebration must never compete with, or crowd out, real help requests. It carries its own gentle limit.

#### 5.9 Disconnection [P1]

- Any connection can be ended, **silently and one-sidedly**. The other person is **never notified**.
- **Chat history is preserved** and remains visible to both — but the thread is **frozen: no new messages**.
- Disconnect sits in the UI alongside **Block** and **Report**.
- **Rationale:** leaving a relationship should never require a confrontation, and preserving history protects both parties if something later needs to be reported. *(Original Q48–Q49.)*

#### 5.10 Block — stronger than disconnect [P1]

Block is a hard separation, not a quiet exit:

- **Mutual invisibility** between the two people.
- **The dispatch engine must NEVER match a blocked pair again — even as strangers.** A blocked person can never be auto-assigned to the blocker's request, appear in a connections wave, or be surfaced in any matching flow.
- **Rationale:** without this rule, someone you blocked could still be auto-assigned to your Urgent request by the matching engine — a serious safety failure.

#### 5.11 Leakage — informal help between connections [P1 decision, P2 feature]

**The situation:** once connected, people can simply message *"can you drop me at the station tomorrow?"* instead of raising a formal request. Every marketplace fears this ("disintermediation"). **For Sapiens it is mostly a good outcome, and it is deliberately not blocked.**

**Why it's acceptable — arguably the mission working:**

- The vision is "a society where helping each other is the default" — **not** "helping each other via an app." Two neighbours who now help each other directly are the stated goal. That is graduation, not churn.
- **There is no revenue to protect.** Uber and Airbnb fight leakage because they take a cut; Sapiens takes nothing. The economic motive to fight it does not exist.
- Forcing friends through a dispatch flow is bureaucratic and cold, and implies the app doesn't trust the relationship it just created.
- It is consistent with 4.6: the system does not police what two people say to each other.

#### What is genuinely lost (accepted costs):

- Invisible helps earn **no Moneta** and don't move the **Goodness Meter** or **Celestial Journey** — a highly active helper could sit at "new moon."
- The badge is also a **safety signal** (1.6); invisible helping makes it systematically understate trustworthy people.
- **No impact data** — weakens grant/CSR conversations (A6) and launch-city decisions (A7).
- No safety scaffolding (meetup code, share-my-meetup, structured dispute window).
- Thinner celebration loop — fewer logged helps, fewer milestones, less community glue.

**The resolution — make the good path the easy path [P1]:** a one-tap "Ask [name] for help" button **inside the connection chat**, which pre-fills a directed request (5.5). If the formal path is one tap away *in the very window where they would have typed the message*, the informal path loses its only advantage — convenience. Most people will take the tap.

**Retroactive help-logging → [P2], deliberately deferred.** A "we helped each other" button (either party taps, other confirms, Moneta awarded) would capture invisible helps — but it is **the single easiest collusion vector in the app**: no meetup code, no physical proof, two taps. Everything else in the system requires a real meeting; this would not. **Decision: measure leakage first.** If most connection-help proves invisible, build logging *with* the pair-repetition anti-abuse layer already in place. **Building the loophole before the lock is backwards.**

---

## BUCKET 6 — INBOX & PERMANENT MESSAGING

The lasting conversation between connections. Covers original questions Q60–Q70.

### 6.1 The frame — two chats, two purposes [P1]

Sapiens has exactly **two** conversation surfaces, and they are deliberately different:

|         | Active-request chat (Bucket 4) | Inbox (this bucket)     |
|---------|--------------------------------|-------------------------|
| Who     | Matched strangers              | Connections only        |
| Purpose | Coordinate one meeting         | An ongoing relationship |

|                  | Active-request chat (Bucket 4)            | Inbox (this bucket)      |
|------------------|-------------------------------------------|--------------------------|
| Lifespan         | Dissolves at completion                   | Persists                 |
| Server-readable  | <b>Always</b> (evidence)                  | v1 yes → <b>[P2]</b> E2E |
| Protection model | <b>Rate limits</b> (can't block mid-help) | <b>Block</b>             |

Everything below follows from that split.

## 6.2 Encryption **[P1]**

**v1 decision: encrypted in transit and at rest, server-readable, with a strict internal-access policy.**

Message content is accessed **only when a report is filed**, and every access is **logged and audited**. Never casual, never browsing.

**Why not end-to-end in v1:**

- True E2E means **Sapiens cannot read anything — including reported messages**. The reporting-and-human-review loop committed to in 1.4 and 4.5 would arrive **empty**, exactly when someone is asking for help. Blocking and reporting inside the inbox (Q62/Q63) would lose their teeth precisely where a relationship has gone bad.
- It collides with **IT Rules traceability obligations** — pre-launch gate **G5**.
- **Honesty over marketing**: it is more on-brand to say "*we protect your messages and can only look if you report*" than to claim a privacy guarantee the product doesn't deliver.

**The privacy policy must state this posture plainly** → ties to **[GATE G3]**.

**[P2] — E2E inbox with client-side reporting.** The endgame: the inbox becomes end-to-end encrypted, and when a user reports, **their own device attaches the relevant decrypted messages to the report**. Privacy by default; evidence when someone asks for help. Harder to build correctly, so deliberately deferred — *ship the reporting loop working, upgrade privacy later*.

**Permanent asymmetry:** the **active-request chat remains server-readable in all phases**. It is the evidence surface for complaints between strangers and cannot be encrypted away. Only the connection inbox is a candidate for E2E.

## 6.3 What the inbox contains **[P1]**

**Connections only.** System and Sapiens notifications live elsewhere in the app.

**Rationale:** the inbox should feel like *people*, not a notification dump.

## 6.4 Disappearing messages **[P1 — no]**

**Not in v1.** Fights the evidence model (4.3/6.2), adds real complexity, and belongs to a different kind of app. → **[LATER]**, if ever.

## 6.5 Identity & naming in the inbox [P1]

- Shown: **first name + photo** — consistent with the disclosure ladder (0.4).
- **Custom nicknames: yes.** A user can rename any connection for themselves ("Raj who fixed the tap"). Genuinely useful and private to the owner.
- **Internal UUID** under the hood.
- **No public handles or usernames.** *Searchable handles would let people find each other directly, which breaks the "no discovery" constitution rule.*

## 6.6 Media & limits [P1]

Same capability set as the active-request chat — **text · photos · voice notes · emoji** — with **slightly higher size caps.**

**Rationale:** consistency is simpler to build and keeps moderation rules uniform across both surfaces. (Voice notes remain an accessibility feature, per 4.4.)

## 6.7 Retention, deletion & media expiry [P1]

- Connection messages **persist indefinitely while connected.**
- After **disconnect**, the thread **freezes but stays visible** to both (5.9). Same for **block**.
- **Delete-for-me: yes.** A user can delete messages or a whole thread from their own view.
- **Delete-for-both: never.** *Removing a message from someone else's device is a manipulation vector* — it would let a bad actor erase evidence of what they said.
- **Media (photos, voice notes) expires after ~6–12 months; text persists.** Cheap on storage, better for privacy, and matches how people actually use old chats. **[OPEN A19]** exact window.
- DPDP requires a user-facing deletion path → **[GATE G3]**.

## 6.8 Presence signals [P1]

- **Delivered / read receipts: at most.**
- **No "last seen." No online status.**

**Rationale:** last-seen and online status create **surveillance and social pressure** — you can see when someone read you and didn't reply. *That is precisely the anxiety machinery of the apps Sapiens exists to reject.* Leaving it out is a quiet statement of what this app is.

## 6.9 Group chats [P1 — no]

**Not in v1.** Group conversations exist **only** for group requests (4.7). Letting connections spin up arbitrary group chats turns Sapiens into a messaging app — scope creep against the constitution. → **[LATER]** if real demand appears.

## 6.10 Account deletion & bans [P1]

When a user deletes their account (a DPDP right, G3) or is banned:

- **The other person keeps their copy** of the conversation. The sender is displayed as "**Former Sapiens.**"
- **Rationale:** the other person's history is *also theirs*, and wiping everything would destroy the evidence trail of a banned user's behaviour.

#### 6.11 Message notifications [P1]

- **Separate from the help-request notification cap** (3.8/A14), with their own gentle limit.
- **Real-time** — a message from a friend should not wait for a digest.

#### 6.12 Anti-harassment rate limits — pre-connection only [P1]

**Rate limits apply only to the active-request chat (i.e. before the two are connected).** The connection inbox relies on **block**, like any messaging app.

**Rationale (the reasoning is the important part):** before connection, a user **cannot block without breaking the help flow** — so a quiet throttle is the protection. After connection, blocking *is* the protection and a throttle would be redundant friction between friends. **[OPEN A20]** exact thresholds.

#### 6.13 Mid-request harassment — the one-action escape hatch [P1]

If someone is harassed **during an active request** — matched, chat open, help not yet done — blocking alone would strand the help. So:

**One action: Cancel + Report + Block.**

1. The match **terminates immediately**; the other party is dropped.
2. The request **re-broadcasts to other helpers**, reusing the existing re-broadcast machinery (0.7).
3. The report goes to **human review with the chat attached as evidence** (possible because the active-request chat is server-readable, 6.2).
4. The pair is **permanently unmatched** by the dispatch engine (5.10).

**Why this shape:** it is the only option where **the victim loses nothing** — the harasser is gone *and* the seeker still gets their help. (Report-only was rejected: it forces someone to keep talking to a person harassing them. Cancel-the-whole-request was rejected: it punishes the victim by taking away their help.)

**Mirror case:** a **helper** harassed by a seeker gets the identical exit — drop the match, report, and **no reliability penalty** for that cancellation. *The no-show logic (0.7) must never punish someone for leaving an abusive situation.*

#### 6.14 Inbox defaults [P1]

Blocking from the inbox = the same hard block as 5.10 (mutual invisibility, **never matched again by dispatch, even as strangers**) · blocked/disconnected threads freeze but stay visible · message search within a thread · standard unread badges.

---

## BUCKET 7 — MONETA, THE METERS & THE CELESTIAL JOURNEY

The reward and reputation layer. Covers original questions Q24–Q25, Q71–Q76, and resolves long-standing open item **A5**.

### 7.1 THE UNIQUE-HELP RULE — the keystone of this bucket **[P1]**

**All rewards and KPIs count only the FIRST-EVER help to a given person.** Helping someone you have already helped earns **no Moneta** and does **not** advance the Goodness Meter, the Celestial Journey, or the leaderboard.

**This is the single most important decision in the reward layer, for three reasons:**

1. **It encodes the mission in the math.** Sapiens exists to *widen the circle* — to make people meet and help as many *different* neighbours as possible. Rewarding only unique helps makes the reward system literally measure the mission.
2. **It structurally solves collusion.** Two friends farming each other now earn exactly **1 Moneta each, ever**. The exploit dies at the *design* level — no detection machinery required. This dramatically simplifies the anti-abuse layer that was parked at P2 (1.5): pair-repetition tracking is largely unnecessary because repetition is already worthless.
3. It keeps one clean definition of "help that counts" across every metric.

**Depth is not thrown away — see 7.2.**

### 7.2 The Steadfast count — recognising depth **[P1]**

The unique-help rule creates a risk: the person who helps the same lonely elderly neighbour every week (groceries, a ride, company) is doing some of the highest community work in the app — yet earns 1 Moneta, once, and their Celestial Journey would stop moving.

**Resolution:** a **separate, visible, honoured counter of total helps** (including repeats) — working name **"Steadfast count"** (*rename freely*). It does **not** drive Moneta or the Meters, but it is **displayed and celebrated** so that *depth* is seen even though *breadth* is rewarded.

- **Breadth** (unique helps) → Moneta, Goodness Meter, Celestial Journey, leaderboard.
- **Depth** (total incl. repeats) → the Steadfast count, its own recognition.

*The app rewards meeting new people, and honours showing up repeatedly for the same person.*

### 7.3 What Moneta redeems for **[P1 decision / phased value]** — resolves A5

- **v1: recognition & access, no material value.** Moneta contributes to **City Saviour eligibility** (as one human-judged input, 7.9) and **community-event access**. Users are told: *help* → *earn Moneta* → *redeem for goodies* (coming), with **NO published exchange rate**.
- **[LATER]: real-world value.** Partner-brand discounts, Sapiens-shop merchandise, free goods — funded by brand collaborations, once they exist.

## Hard rules (all phases):

- **Moneta must NEVER buy priority in dispatch or better access to help.** That would recreate the money economy inside the app (paying to be helped first) — a constitution violation.
- **Moneta is permanently non-transferable (7.5).**

### 7.4 ⚠ The redemption ↔ anti-abuse dependency [GATE]

**Real-world redemption must NOT launch before the anti-collusion layer exists.**

**Why:** while Moneta buys nothing, farming is pointless — which is exactly why deferring anti-abuse to P2 is safe today. **The day Moneta buys real value, farming becomes profitable,** and every deferred defense becomes load-bearing. The unique-help rule (7.1) already removes the *simplest* farming path, but sophisticated abuse (many fake accounts, each helped once) would still pay.

**Sequence rule:** ship recognition-only Moneta now → build the anti-abuse layer → *then* enable real-world redemption. **Never enable redemption first.** Added to the pre-launch gate cluster (see G7, 2.4).

### 7.5 Moneta value, transferability, expiry [P1]

- **No published cash value in v1.** The ₹10-per-Moneta figure is an **INTERNAL planning number only** — used later to negotiate partner rewards ("a 500-Moneta reward should feel like ~₹5,000 of value to us"). **Never shown to users as an exchange rate. Rationale:** publishing "1 help = ₹10" converts Moneta into a monetary instrument — helping-for-money in disguise — which breaks constitution rule 1 and triggers RBI prepaid-instrument / voucher / tax exposure. Keeping value implicit ("500 Moneta unlocks this") preserves the constitution while allowing partner economics. → **regulatory review is a gate item (G7).**
- **Non-transferable, permanently.** Users cannot send Moneta to each other, in any phase. *Transferability would make Moneta a currency and spawn a grey market ("₹200 for 50 Moneta") once it has value.*
- **Never expires.** Expiry punishes people for not spending — hostile to a goodness token.

### 7.6 Ledger architecture [P1]

**Two separate numbers, always:**

- **Lifetime Moneta earned** — never decreases. Drives the Celestial Journey, milestones, and (via unique helps) the Goodness Meter.
- **Spendable balance** — earned minus spent. Only becomes meaningful once redemption exists.

Plus an **immutable, append-only ledger:** every Moneta event recorded with help ID, timestamp, and both parties — required for dispute reversals (0.8 clawback) and fraud investigation.

**Wallet:** a simple view showing lifetime earned + spendable balance + history. **No cap** (capping punishes the best helpers; if hoarding ever becomes a problem the answer is decay/expiry, not a ceiling — and neither applies in v1).

## 7.7 The three Meters [P1]

Each abstract meter maps onto a concrete mechanic already in the system. **Three meters, each meaning one distinct thing** — no fourth gauge.

| Meter           | Mechanic                            | Scale       | Measures                                                                              | Visibility                 |
|-----------------|-------------------------------------|-------------|---------------------------------------------------------------------------------------|----------------------------|
| Trust Meter     | Average star rating from feedback   | 0–5 stars   | <i>Who you are to people who met you.</i> Reflects <b>recent</b> behaviour naturally. | (see 7.8)                  |
| Goodness Meter  | Unique helps, car-style speedometer | 0–100 index | <i>How wide is your circle.</i> <b>Lifetime climb.</b>                                | <b>Owner-facing</b>        |
| Spiritual Meter | The Celestial Journey (moon → sun)  | Stages      | The emotional/spiritual heart; the <b>public trust signal.</b>                        | <b>Public</b> (stage only) |

- **Goodness Meter curve:** a responsive diminishing-returns curve,  $100 \times (1 - e^{-(\text{uniqueHelps} / k)})$ , tuned so a normal helper sees *the needle move* ( $\approx 100$  unique helps → a visibly meaningful score) while 100 is approached only by the exceptional. [OPEN A21] tune  $k$ .
  - *Rejected:* literal  $\text{uniqueHelps} \div 10$  — it would sit near zero for years (100 helps = 10/100) and a speedometer that never leaves zero isn't motivating; a car analog implies movement you can feel.
- **Goodness Meter does NOT decay** — it is a pure lifetime climb (unique helps never decrease). **The Trust Meter carries "latently"** (a quiet or bad patch shows in stars). Each meter thus means one clean, non-overlapping thing.
- **Happiness** is the *feeling/tagline* of the Goodness Meter ("spreading happiness to more people"), **not a fourth meter.**

## 7.8 The Celestial Journey in-app [P1]

- Thresholds (unchanged, now meaning **unique people**): crescent **10** · half moon **50** · full moon **100** · sunrise **500** · golden sun **1000** → light joins the shared galaxy.
  - *Accepted consequence:* golden sun now = **1000 different people** — a near-mythic lifetime achievement, which makes the badge genuinely rare and meaningful. Milestones kept as-is by decision.
- **Badge shows the STAGE publicly; the exact number is owner-only.** Keeps it a *trust signal*, not a public scoreboard.
- **The shared galaxy is symbolic and beautiful in v1;** a real live galaxy of actual users' suns is a [P2] visualization.



## 7.9 Leaderboard [P1]

- **Ranked by: new lifetime-unique people met within the time window** (Model A). Meeting a person for the first time ranks you *that month*; helping a repeat person ranks you nowhere (they aren't new). Measures "*how many new people did you reach this month.*"
  - **Rationale:** perfectly consistent with the unique-help rule, impossible to game by repeat-helping, and purest to the mission (widen the circle). *Accepted tradeoff:* someone whose helping is deep-not-broad ranks low — but their depth is honoured via the Steadfast count (7.2).
- **Automatically opt-in** (*as written — every eligible helper appears; revisit if quiet helpers object* → A22).
- **Filters:** Zip code / City / State / Country.
- **Reset:** monthly.
- **Prize in v1: recognition only** — a mention, a badge, a moment. Cash rewards and roles are [LATER].
- **Not raw help count:** ranking by uniques (not total) rewards breadth and resists gaming.

## 7.10 City Saviour selection — human-judged [P1 principle / LATER operationally]

City Saviours are selected by **community nomination + interview**; the leaderboard is **one input, not the selector**.

**Rationale — this is what keeps the leaderboard warm:** if topping the board *got you hired*, the board becomes an economic competition and gaming becomes rational. Making selection human-judged (and saying so publicly — "*topping the board doesn't get you hired; being genuinely good does*") defuses the gaming incentive before it forms. Operational design is [LATER]; the *principle* is locked now because it defines what the leaderboard is.

---

# BUCKET 8 — COMMUNITY MOMENTS (the "feed" question)

The social/celebration surface. Covers original questions Q77–Q93. Resolves the app's hardest philosophical tension.

## 8.1 The reframe — the job is not "build a feed" [P1]

The real needs are **(1) prove the app is alive** (so a newcomer trusts it) and **(2) pull people back**. Every social app solves these with an infinite feed — which is precisely the dopamine trap the Sapiens constitution rejects ("we don't farm attention," "no feed to scroll").

**Governing principle for this entire bucket:** *the thing that brings someone back should be a real event involving them — not an infinite stream of strangers to consume*. Keep the **jobs**, refuse the **mechanic**.

## 8.2 The surface — finite "community moments," not a feed [P1]

A **calm, finite "community moments" surface**: real completed-help moments (mutual-5-star selfies, milestones). **It has a bottom**. Chronological or lightly curated — never algorithmic, never infinite.

- *Rejected*: a real social feed (the thing the thesis rejects) and "no shared surface at all" (leaves a brand-new user with nothing to see, failing the newcomer need).

## 8.3 Newcomer vs. retention — solved SEPARATELY [P1]

The two needs are different and get different answers:

- **Newcomer** ("show me this is real"): a curated **"This is Sapiens" welcome view** — real anonymized moments, the activity map (8.6), the feeling of the place — shown *before* the person has any activity of their own.
- **Active user** ("bring me back"): driven by **personal notifications** — a nearby request they can help with, a connection's milestone, their request being accepted — **not** by a feed.

**The healthy inverse of engagement apps**: the shared surface matters *most before you are active* and *fades* once you are. The app does not need to hold your attention; it needs to get you to help someone and then let you go.

## 8.4 The mutual-5-star selfie [P1]

The signature social object — a celebration of a *moment of help*, not a profile-builder.

- **Double opt-in, always**. Posting a photo of two people to a community requires **both** to actively agree — a **separate action from giving 5 stars**. (*Privacy hard-line.*)
- **Revocable unilaterally, anytime**. Either party can remove the selfie afterwards, no questions asked. *A permanent public photo you can't retract is a serious privacy problem.*
- **Visibility = public, but privacy-shaped + area-scoped (B + C)**:
  - The moment celebrates the **deed and the feeling** ("two neighbours, a house moved in Koregaon Park"), with faces shown **only because both consented**.
  - **NOT a searchable directory**. **NO tap-through to a browsable profile**. **NO location precision**. **Not searchable by person**.
  - **Area-scoped** — you see your locality's moments, not a global scroll. It feels like *your community's good news*, not a global content stream.
- **The two-layer defense (documented so the reasoning is airtight)**:
  1. **In-app misuse is blocked by connection-gating** — surfing the surface cannot lead to *connecting* with anyone (no help → no connection, locked in Bucket 5), and only the same essential info a seeker already sees is ever exposed.
  2. **Off-platform exposure is limited by B+C scoping** — a public *global* face-feed could be screenshotted and a person identified/found elsewhere; area-scoping + no-searchable-directory shrinks that surface hard. *The defense is both together — gating stops in-app misuse, scoping limits off-platform exposure — not either alone.*

## 8.5 Milestones in the surface [P1]

Milestones (10/50/100/500/1000 unique helps) **appear in the community moments surface** — because a real person reaching a milestone has *real motivating impact* on others, more than a bare number.

- Still also delivered as a notification to the person's connections (5.8).
- **Never rendered as a leaderboard-style flex** — it's a celebration, not a ranking.

## 8.6 The activity map / counter — aggregate only [P1]

A **live, anonymized activity map/counter** — "37 helps near you this week" — alongside the moments surface. Proves life, builds local pride, and helps the cold-start-of-trust problem (A3).

- **Aggregate only.** Heat/counts by area, **never an individual locatable event.**
-  **Safety hard-line: never plot an individual live help on a public map — especially never an SOS.** A precise pin can reveal where a vulnerable person is (a "help at home" or SOS pin = a person alone at a known spot). **Counts and warmth, not coordinates.**

## 8.7 Interactions — one gesture, no scoreboard [P1]

- **The only interaction is  "Appreciate."**
- **No comments in v1** — comments invite negativity, moderation load, and social-performance dynamics the app avoids. → [LATER] only if ever justified.
- **The  count is hidden / soft** — the sender feels the warmth; there is **no public tally.** *Even a heart becomes a vanity metric if the number is public ("200 hearts vs. my 3").* Keep the gesture, kill the scoreboard — consistent with rejecting vanity metrics everywhere (5.7, 7.8).

## 8.8 Content boundary — organic only [P1]

- The in-app moments surface carries **organic user moments only** (selfies, milestones).
- **Curated real-life impact stories and interviews stay on the website** (its Phase-2 content), not in-app. *Keeps the app surface authentic and low-maintenance.*

## 8.9 Anti-addiction — HARD RULES (permanent, all phases) [P1]

These are **banned outright, regardless of any future engagement or growth pressure.** They exist so no later optimization can quietly re-introduce the dopamine trap:

1. **No infinite scroll.** The surface has a bottom.
2. **At the bottom:** *"You're all caught up — go help someone,"* and it **stops.** No "load more," no refill, no pull-to-refresh-for-more. **The bottom is the feature.**
3. **No autoplay.**
4. **No public like/heart counts** (8.7).
5. **No streaks that punish absence.**
6. **No notification bait / engagement-farming pushes.**
7. **No algorithmic outrage- or engagement-optimization.** Ordering is chronological or lightly, transparently curated — never attention-maximizing.

These bans are constitutional (rule 6, "we don't farm attention") rendered concrete.

---

## BUCKET 9 — PROFILE ARCHITECTURE

Largely assembly — earlier buckets already decided what a profile *contains*. This bucket confirms layout, the owner-vs-visitor split, editing, and sharing. Covers original questions Q17–Q18, Q94–Q104.

### 9.1 The frame — two profile views [P1]

Sapiens has **two radically different profile views**, and almost every decision here is "which view?":

- **Owner view** — rich, private, your whole journey.
- **Seeker-at-match view** — minimal, safety-essential only, governed by the disclosure ladder (0.4).

### 9.2 The TWO-PHOTO model — the key decision [P1]

One photo was doing two different jobs; they are now split:

|                 | KYC / verification selfie                       | Display / profile photo                |
|-----------------|-------------------------------------------------|----------------------------------------|
| Job             | Proof you are a real, specific human            | Recognition ("is this who showed up?") |
| Editable?       | <b>Never</b>                                    | <b>Yes</b> — user can update it freely |
| Shown publicly? | <b>Never</b> — lives in the verification system | Yes — seen at match + on profile       |

**The rule: your display photo can be anything you like, as long as it is verifiably still you.**

- **Every display photo is auto face-matched to the KYC selfie before it goes live** (the dating-app model — the "verified" status holds because every photo is matched to the verified one). (*Face-match method A.*)
- **On match-fail or uncertainty: reject the new photo and keep the last verified-matching one;** let the user try again. **Repeated failures are soft-flagged to Trust & Safety** — a single failure is innocent (lighting, sunglasses), but a *pattern* of trying to slip in a mismatched face is a signal.
- **At the meetup, identity is confirmed by the DISPLAY photo** (what the seeker has been looking at) — trustworthy precisely because it is face-matched to KYC. The KYC selfie stays private.
- *Rejected: free editing with no check* — a user could set someone else's photo, collapsing meetup photo-verification and the entire identity-only safety model (1.1/2.2).

**Why this matters:** it gives users the expressive, "a picture says more than words" profile they want **without** weakening the verified-identity safety model. Photo expressiveness and identity proof are now decoupled and each done properly.

### 9.3 Owner profile — contents [P1]

Everything, with **exact numbers owner-only** (7.8):

- **Celestial badge + stage** · the **three meters** (Trust 0–5 / Goodness 0–100 / Spiritual = the Journey)
- **Lifetime Moneta earned + spendable balance** (7.6)
- **Unique-help count** + the **Steadfast (repeat-help) count** (7.2)
- **Connections** list (visible to owner + connections, 5.7)
- **Help history** — owner-only (9.5)
- **"Ways I help"** — the categories you've opted into, shown as a small identity statement ("I help with rides, tutoring, moving"). **Editable** (9.6).
- **Impact framing over raw numbers** — "You've reached 40 neighbours" rather than a bare "40." Reinforces the mission; reads warmer.
- **A private "your journey" timeline (owner-only)** — first help, milestones, the moon's progress. A quiet personal keepsake, **not a public flex** — the kind of thing that brings people back for *healthy* reasons.

**Deliberately excluded:** anything that becomes a public number or comparison surface — no "top X%," no scoreboard framing. (*Consistent with the vanity-metric rejection across 5.7, 7.8, 8.7.*)

### 9.4 Seeker-at-match view — the safety-critical view [P1]

At "hand raised" (0.4), the seeker sees **exactly**:

- **First name** · **verified display photo** · **distance** · **Celestial stage** · **star rating (Trust Meter)** · **"member since."**

**And nothing more until connection.** *The seeker needs enough to feel safe, not a dossier.* (Star rating added as the most decision-useful trust signal; "member since" as a light tenure cue.) Everything else — meters detail, counts, history, full profile — unlocks only on **Connection** (5.4).

### 9.5 Help history — owner-only, relationship-shaped [P1]

- **Your connections appear named. Non-connection helps are anonymized / shown as counts** — "you've helped 40 people," but you do **not** keep a browsable file on strangers you met once.
- **History is owner-only** — never visible to anyone else. A *visible "who they helped" list is a surveillance surface* and is forbidden.

### 9.6 Editing — what's editable vs. locked [P1]

- **Editable:** short bio/intro · **Ways I help** · availability / quiet hours · radius maximum · notification preferences · trusted contacts · **display photo** (face-matched, 9.2).
- **Locked (would let someone misrepresent a verified identity or fake reputation):** legal name · KYC identity · **KYC selfie** · Celestial badge · meters · Moneta · help history.

The "missing category" suggestion box [P1]

When choosing "Ways I help," if a category is missing, the user can **type it and submit it for review**. It is **not added live** — it goes to the team, and on approval **enters the taxonomy config** (Bucket 3's editable-data model).

**Why this is valuable:** it is the **input pipe that keeps the taxonomy alive** (3.11) — no app update needed to add categories — *and* doubles as **demand research** (the most-requested categories signal where to expand). The **review gate is essential** — it stops the category list filling with junk or abuse.

9.7 Profile sharing — none in v1 [P1]

**No public, shareable profile.** *A shareable, viewable profile is a findable person — it would quietly reintroduce the discovery/surfing the constitution bans.* Users share an **invite to Sapiens**, not a profile.

9.8 Requests from the profile [P1]

- **A directed request can be launched from a connection's profile** (contextual — you're asking *this* person).
- **General requests belong on the home surface**, not a profile. (*Low-stakes UX; original Q104.*)

9.9 Connections list & mutual profiles [P1]

- Connection count + list visible to **owner + their connections** (5.7).
- Tapping a connection opens their **fuller profile** and the shared history between you — *connection is the thing that unlocks fuller visibility* (5.4).

---

## BUCKET 10 — UI/UX & NAVIGATION

The original Module 1, designed last on purpose — behaviour was decided first, so most screens now follow from it. Covers original questions Q1–Q6, Q9, Q113–Q118.

10.1 The UX thesis — the 70-year-old test [P1]

Every screen must pass one test: **can a frightened, non-technical 70-year-old use it in one try** (the woman "with no technical knowhow," the unwell elderly man — the app's own core use cases). **Radical simplicity** is the design philosophy: the home screen answers "*what do I do?*" in one second; everything else gets out of the way. This is the rare app where doing *less* on screen is the feature.

10.2 Design language & typography [P1]

- **Inherits the website's soul:** warm paper / ink / spark-orange + the celestial motif. Continuity between site and app.
- **Nunito Sans is the workhorse** — all functional/body/UI text. Chosen over Space Grotesk because it is warmer and more legible at small sizes and for older eyes (the accessibility bar); Space Grotesk read cooler / more "tech product." **Cabin Sketch** is reserved for celebratory brand

moments only (milestones, the Journey). [OPEN A23] final font sign-off (Space Grotesk held as a headings-only option if a cooler, more modern app identity is later wanted — never for body).

- Generally a "day" feel, with the celestial **night** reserved for the Journey.

### 10.3 Navigation skeleton [P1]

- **Bottom tab bar** (familiar, thumb-reachable, and older users already know the pattern from WhatsApp).
- **Four tabs: Home · Moments · Inbox · You.**
- **SOS is a persistent, guarded button — not a tab** (10.8).
- The **two core actions live on Home** (10.4).

### 10.4 The home screen [P1]

Calm and uncluttered. Contains:

- A greeting + the **guarded SOS button** (top corner).
- **Celestial Journey glance** in *impact language* — "You've reached 40 neighbours," not "40" (7.7, 9.3).
- **Aggregate activity counter** — "37 helps near you this week" (proves the place is alive; 8.6).
- **The two big actions: "I need help"** (primary) and **"Help someone"** (opens the bounded Help-now list, 10.6).

### 10.5 The raise-help flow (the seeker's journey) [P1]

**Stripped to three taps to dispatch:** Tap "I need help" → pick a category → say what & when → send.

Every extra step loses a distressed user.

- **Category selection = big visual tiles** (icon + plain words) for the 8 parents (3.11) — not a small text menu. Accessibility-first.
- **Describing the need:** a **dropdown/tile selection in MVP**, plus a short optional description field.
  - [P2] **voice-note input with sentiment/category analysis** — the AI **suggests** a category for the user to **confirm** (never silently auto-assigns), and **must never auto-trigger SOS** from interpreted speech (a mis-read emergency is a safety failure). Needs Indian-language/accent handling (ties to A8).
- **Urgency (resolves A1): pre-filled from the category's typicalUrgency default** (3.11) — a ride = Everyday, "stranded" = Urgent — and the user may *nudge* it with a light control. **SOS is never in this control** (it's the separate guarded button). System-defaults + limited self-escalation.
- **The waiting screen: a calm, reassuring live status** — "Finding someone nearby... 3 neighbours notified" — **never a dead spinner**, with **widen-search** and **SOS** reachable, and honest "notify me when someone can help" messaging if nobody's available (3.10). Tone: calm and reassuring, never alarming.

### 10.6 The helper's side [P1]

- **Being pinged:** the helper sees **help type · rough distance · approximate area · urgency · how long ago** — enough to decide, but **not** the seeker's identity or exact location (disclosure ladder, 0.4). **One-tap "I'll help"** vs. ignore.

- The "Help now" bounded list (from Home): **finite, matched to the helper's categories, nearby, with a bottom** — never an infinite feed. **Ordered urgent-first, then nearest** (someone actively looking to help should see who needs it most). Shows the same limited pre-accept info as a ping.
- **After "I'll help" (waiting for confirm/veto)**: in first-to-accept tiers the helper matches instantly; in Urgent auto-assign / veto flows they may wait. **If the seeker vetoes them: a gentle "this request was matched with someone else"** — the same **silent, no-rejection-sting** principle as connection declines (5.3). A helper must never feel personally refused.
- **The active-help screen (helper view)**: mirrors the seeker's tracking screen + **"I'm on my way / I've arrived"** status buttons that drive the no-show machinery (0.7) + chat + **meetup-code entry** + **"mark as done."** The **one-action Cancel + Report + Block** (6.13) is reachable here.
- **Helper controls** (categories, radius max, quiet hours) live in **Profile** → **"Ways I help"** (9.6) — never cluttering Home.

## 10.7 Onboarding / first-run experience [P1]

**One combined first-run flow** (three jobs, one journey):

Sign up (email + phone, OTP) → short walkthrough → set trusted contacts → "This is Sapiens" welcome view → free to look around → KYC gate the moment they try to help or ask for help.

- The walkthrough teaches *mechanics* (the two actions, the Journey, SOS) — **short, teach-by-doing, a few warm screens**, skippable. Distinct from the **"This is Sapiens" welcome view** (Bucket 8), which builds *belief*. Both greet a newcomer; different jobs.
-  **SOS must be self-explanatory even if the walkthrough is skipped** — on trigger, a confirm screen states plainly *"This alerts your trusted contacts and lets you call 112."* People skip tours; the safety-critical part cannot depend on the tour.
- **Onboarding is where the 1–3 trusted contacts are captured** (SOS Layer 1 dependency, 1.2). **Strongly prompted but skippable, then re-prompted before it matters; required before a user is fully active.**
- **Trusted contacts, walkthrough, and welcome view are one flow.**

## 10.8 Unverified-user access [P1]

An unverified (pre-KYC) user can **browse the welcome view and Moments, and set their preferences** — enough to build belief. The instant they tap **"I need help"** or try to **respond to someone**, the **KYC gate** appears. Nothing that requires trust is available before verification (2.1).

## 10.9 SOS access & the soft daily limit [P1]

- **Persistent guarded button** (top corner of Home), activated by **press-and-hold or slide-to-confirm** (the iPhone Emergency SOS pattern people already know) — always reachable, never a mis-tap.
- **Soft daily limit (2–3), NEVER a hard lockout.** The critical shape:
  - After the limit, the app **does not remove the SOS button**. Further presses **still route to one-tap 112 + trusted-contact alerting** (the life-safety floor), but **stop generating fresh community/system SOS events**, and the account is **flagged to Trust & Safety**.



- **Rationale:** a person can genuinely have a second real emergency; a hard cap would block a real cry for help — catastrophically worse than SOS spam. The limit caps the *abusable* part (new blasts), never the *call-for-help* part (112 + contacts). *No emergency system locks the button; you handle abuse by accountability, not lockout.*
- **Accountability:** every SOS has an "are you safe now?" resolution step; repeated/false SOS → escalating friction, warnings, ultimately a ban. Combined with the soft limit, this contains spam without endangering anyone.

## 10.10 Location & maps — no Map API in v1 [P1]

**"Tracking" is two separable things, and only one costs money:** the *location data* comes from the phone's built-in **GPS (free)**; only *drawing a rendered map* incurs Maps-API charges. So v1 uses the free part without the expensive part.

### MVP approach (zero map spend):

1. **Navigation = deep-link out to the user's own maps app.** Meeting-point **coordinates** are handed off via a universal maps link; Google/Apple Maps opens with the destination set and does turn-by-turn. Their app, their (zero) cost, better navigation than we'd build.
2. **Passing the seeker's location to the helper is just passing coordinates** (a lat/long pair) phone-to-phone through the backend — **no map involved, free.** A **"Navigate"** button on the helper's active screen deep-links those coordinates into their maps app.
3. **"Where's my helper" = reassuring text/ETA, not a live dot.** The helper's phone reports its GPS (free); the seeker sees *"Raj is 400m away · ~3 min,"* updating as they approach, then *"Raj has arrived."* For a scared, low-tech user this is arguably **calmer and clearer than a moving dot.**

**Staged-disclosure rule preserved:** the **precise** meeting location (and the navigate deep-link) is generated **only at the moment of confirmation** (0.4/0.5) — never during matching.

### Pluggable architecture (so an embedded map slots in later with no rework):

- All location data flows through **one "location provider" module** — GPS in, coordinates + distance + ETA out. The rest of the app only talks to *that*, never a map SDK directly.
- The **display is a swappable component:** v1 renders the text/ETA card; a future version renders a live map in the **same slot**, fed by the **same** service.
- **"Show a location"** is an abstraction: v1 opens the external maps app; a later version draws an embedded map — same call, different implementation.
- **Not locked to Google:** when embedding later, **Mapbox / MapLibre (open-source) / OpenStreetMap** are far cheaper or free at scale — the abstraction means picking the cheapest map is a display swap.

**Accepted limitation (conscious):** until a map is embedded, the seeker sees distance + ETA + "arrived," **not** the helper's live path/dot. Acceptable for the calm, low-tech-friendly design; upgradeable later.

---

### 10.11 Profile screen — the "You" tab [P1]

Contents were decided in 9.3; this is arrangement, top to bottom: **Display photo + name + Celestial badge** → **the three meters (Trust / Goodness / Spiritual)** → **impact numbers (neighbours reached · Steadfast count · Moneta)** → **"Ways I help" (editable)** → **private journey timeline** → **connections** → **settings**.

### 10.12 Notifications surface [P1]

- A **dedicated notifications view behind a bell icon on Home** — **not** a fifth tab (the four-tab skeleton holds, 10.3).
- Help-request pings, connection milestones, and system messages gather there.
- **Message** notifications route to the **Inbox**; **help pings** also surface as the **Help-now list** (10.6).
- All notification budgets/caps from earlier buckets apply (3.8, 5.8, 6.11).

### 10.13 Empty states [P1]

**Every empty state is warm and directive, never blank** — on day one in a new area, screens will be empty, and a dead-looking app loses the user. Empty states are gentle invitations to act:

- Empty Moments → *"No moments near you yet — be the first to help someone."*
- Empty connections → *"Help someone, and your first connection begins here."*
- Empty inbox → similar warm prompt.

---

## BUCKET 11 — TECH STACK & INFRASTRUCTURE

How Sapiens is built. Determined by the product above. Covers original questions Q7–Q8 and the Security section (Q119–Q121). **The website's stack reaches far into the app** — same React, same Supabase, same PostHog, same Vercel-hosted admin pattern — so the solo-buildability that shipped the site carries over.

### 11.1 App platform — React Native + Expo [P1]

**Cross-platform, one codebase, both stores** (native Swift + Kotlin is untenable solo). Within that: **React Native + Expo. Rationale:** reuses the **React** already used for the website (same language, components, mental model — not starting from zero); **Expo removes almost all solo-mobile pain** (build pipeline, push, GPS, the camera for the KYC selfie, over-the-air updates); and **Claude Code is fluent in it.** (*Flutter considered and rejected — excellent, but a new language (Dart) that discards the React synergy.*) **iOS + Android both at launch** (cross-platform makes both nearly free).

### 11.2 Backend — Supabase [P1]

**One backend for site and app**, already known. Supabase provides, in one place:

- **Postgres** — the data model.

- **Realtime** — live dispatch pings + chat (to *open* apps).
- **Auth** — email + phone OTP.
- **Storage** — photos + voice notes (with retention rules 4.3, 6.7).
- **Row-Level Security (RLS)** — the privacy rules designed across all eleven buckets map **directly** onto RLS policies.
- **Edge Functions** — server logic (dispatch engine, etc.).
- **PostGIS** — the geospatial "who's nearby" queries, free and built in.

### 11.3 The dispatch engine — the custom core [P1]

The **one genuinely hard, custom piece**: wave-based, geospatial, timeout-driven matching (find nearby opted-in helpers → ping in waves → timeout → re-broadcast). No backend gives this for free. Shape:

- **PostGIS** for the radius + category eligibility query (3.2, 3.3).
- **Edge Functions + scheduled jobs** for wave and timeout logic (0.7, 3.3).
- **Realtime + Push** to deliver pings (11.5). **Built incrementally**, but the core objective (reliable dispatch) must be achieved. *This is the part most likely to test "solo + Claude Code" — the geospatial wave/timeout logic is novel enough to need careful iteration and seeded-data testing. Budget the most effort here.* May grow into a small dedicated service later.

### 11.4 Dispatch timing mechanism [P1]

- **v1 (A): Supabase scheduled functions / pg\_cron** — native, simplest, runs on an interval (checks ~every minute), so timing is **approximate** ("roughly 2 minutes" is acceptable for dispatch).
- **Later (B): a dedicated queue/worker** — precise, scales better — if precision or scale demands it. **A now, B later** — the "grow into a service" path.

### 11.5 Realtime vs Push — the delivery split [P1]

- **Supabase Realtime** → live updates to an **open** app.
- **Expo Push** (system notification) → reaching a **backgrounded/closed** app.
- **A dispatch ping fires BOTH**. Critical, because most helpers won't have the app open when a request arrives.

### 11.6 Supporting services [P1]

- **Push: Expo Push** (wraps Apple/Google; free, simple).
- **Auth**: Supabase Auth (email + phone OTP).
- **Media**: Supabase Storage + expiry rules (4.3, 6.7).
- **Location: Expo Location** (free GPS) + deep-link to external maps (10.10). No Maps API in v1.
- **Chat/realtime**: Supabase Realtime for both the active-request chat and the inbox in MVP.
- **Analytics: PostHog** (already used on the website).

### 11.7 SMS [P1]

Trusted-contact SOS alerts (1.2) and phone-OTP both send SMS — a small but real running cost.

- **Route all SMS through one "SMS provider" module**, so **MSG91 / SmartPing / Twilio** (or any Indian gateway) is a swap, not a rebuild.

-  **India DLT registration:** sender IDs + templates must be pre-registered on the government-linked DLT platform before sending transactional SMS (OTP, SOS). A real onboarding step with any Indian provider → **pre-launch checklist**.
- **[OPEN A24]** choose SMS provider (needs live research on current pricing/deliverability).

### 11.8 Offline / degradation [P1]

India means patchy networks, and this app has safety-critical moments.

- **SOS degrades gracefully:** with poor/no data it still fires **trusted-contact alerts by SMS** (Layer 1 was always SMS-based, 1.2) and offers the **112 dial** (works with zero data).
- Ordinary features simply show "you're offline."

### 11.9 Admin / Trust & Safety dashboard [P1]

**Not optional — the safety model has no operator without it.** A simple Next.js web dashboard on **Supabase** (mirrors the website admin, Vercel-hosted), MVP-scoped to:

- View/resolve **reports**.
- Read **flagged active-request chats** as evidence (possible because that chat is server-readable, 6.2).
- **Ban/suspend** users.
- **Approve suggested categories** (9.6). Powers the reporting loop (1.4, 6.13) and the SOS accountability flow (10.9).

### 11.10 App Store review — a pre-launch reality [GATE-like]

Apple (and Google) review **safety apps and user-generated-content apps strictly** — an app where strangers meet in person, with SOS and user content, will face scrutiny, and reviewers want to see **moderation, reporting, and safety mechanisms working** before approval. *Good news: all of it is designed.* But the review itself is a real gate alongside the legal gate → **pre-launch checklist**.

### 11.11 Cost posture [P1]

- **Free tiers for MVP + initial testing; paid only at real scale.**
- **The only unavoidable paid services at MVP:** the **KYC provider** (per-verification, 2.2) and **SMS** (per-message, 11.7). Everything else (Supabase, Expo, PostHog, Vercel admin) runs free until scale.

### 11.12 Security consolidation (Q119–Q121) [P1]

Mostly **already designed across the buckets** — consolidated, not reopened:

- **RLS** enforces every privacy rule (11.2).
- **Encryption in transit and at rest;** message content accessed **only on report, logged and audited** (6.2). E2E inbox is [P2]; the active-request chat is permanently server-readable (evidence).
- **The unique-help rule structurally dissolves collusion** (7.1) — the strongest anti-abuse mechanism in the system, requiring no detection machinery.

- **Rate limits** on pre-connection chat (6.12); **KYC-verified, liveness-matched identity** as the accountability backbone (1.1, 2.2); **soft SOS limit + accountability** (10.9); **blocked pairs never re-matched** (5.10).
- **Minimal sensitive-data storage** — never a raw Aadhaar number (2.2); minimal retention windows (4.3, 6.7).
- Deferred anti-abuse (pair-repetition, fake-account farming) is safe to defer **only while Moneta has no cash value** — the redemption↔anti-abuse gate (7.4/G7) enforces the order.

### 11.13 The full stack, at a glance

| Layer           | Choice                                                                           |
|-----------------|----------------------------------------------------------------------------------|
| App             | React Native + Expo (iOS + Android, one codebase)                                |
| Backend         | Supabase (Postgres · Realtime · Auth · Storage · RLS · Edge Functions · PostGIS) |
| Dispatch engine | PostGIS query + Edge Functions + <code>pg_cron</code> (→ dedicated worker later) |
| Live delivery   | Supabase Realtime (open app) + Expo Push (closed app) — both fire                |
| Location        | Expo Location (free GPS) + external-maps deep-link (no Map API v1)               |
| Media           | Supabase Storage                                                                 |
| SMS             | Provider module (MSG91 / SmartPing / Twilio) — DLT-registered                    |
| Admin / T&S     | Next.js web dashboard on Supabase (Vercel)                                       |
| Analytics       | PostHog                                                                          |
| Build           | Solo, with Claude Code                                                           |
| Cost            | Free tiers → paid only at scale (KYC + SMS the only MVP paid services)           |

## 12. STRATEGY — LAUNCH, METRICS & FUNDING

12.1 Launch strategy — closed communities first **[P1]** — resolves A7

**Launch through closed, connected communities: colleges, companies, organizations, gated housing/societies** — not open cities. **Rationale (this single choice solves three problems at once):**

- **Density** — a hyperlocal app needs people nearby; a shared campus/company *is* a dense pocket. (An empty hyperlocal app feels dead and loses users.)
- **Cold-start** — a bounded community reaches critical mass far faster than a whole city.
- **Baseline trust** — people in a shared campus or company already have a **thread of accountability** between them *before* Sapiens layers verified identity on top, making the identity-only safety model (1.1) meaningfully safer at launch. This matches the original vision's "first users" (universities, companies, mega housing complexes) and the "planted Sapiens" seeding idea.

12.2 Success metrics **[P1]** — resolves A11

- **North-star: completed unique helps** — the mission itself, made measurable (breadth of real help delivered).
- **Guardrail 1: safety incidents** — inverse metric; **must stay near zero**. The one number that can end the product.
- **Guardrail 2: helper retention** — are people *staying willing* to help? (Supply is the fragile side of the marketplace.)

12.3 Funding **[PARKED — dedicated session]** — A6 remains open, consciously

The organization has **no revenue by design**, but running it is not free (SMS, KYC per-verification, maps/servers at scale, eventual City Saviour stipends). **"No money in the app" ≠ "no money to run it."** Candidate paths to explore in a dedicated session: **grants · India's mandated corporate CSR funds (a strong fit) · donations · impact investment · the Premium Choice upgrade (0.9)**. *Deferred deliberately — not forgotten.*

---

## 13. MVP SCOPE — P1 AT A GLANCE

The boundary the Build Specification will implement. **In P1** (the first shippable app):

- **The spine:** raise request (category + urgency + Now/Scheduled) → dispatch (implicit availability, adaptive radius, waves, category-filtered) → match (first-accept / Urgent auto-assign + seeker veto) → confirm (staged disclosure, precise location released, chat opens) → meet (public-point default, meetup code) → complete (helper done → seeker confirm/auto-confirm) → Moneta (flat 1, **unique helps only**) → double-blind rating → optional mutual Connection.

- **Safety:** identity-only model · SOS Layers 1+2 (trusted contacts by SMS + one-tap 112) · everyday-safety scaffolding · share-my-meetup · prefer-women-helpers · in-app reporting + the T&S dashboard · the one-action Cancel+Report+Block.
- **Identity:** KYC via provider (Aadhaar-primary + DL fallback + liveness), minimal storage.
- **Comms:** active-request chat (text/photo/voice/emoji, server-readable, dissolves at completion) · connection inbox (encrypted in transit/at rest, block-based).
- **Reward/reputation:** Moneta (recognition-only value) · Goodness/Trust/Spiritual meters · Celestial Journey · Steadfast count · monthly unique-helps leaderboard (auto-opt-in).
- **Social:** finite community-moments surface · double-opt-in selfies (area-scoped) · milestones · aggregate activity map · ❤️ appreciate (no counts) · the anti-addiction hard rules.
- **Profile & UI:** two-photo model (KYC vs editable display, face-matched) · four-tab navigation · the raise-help and helper flows · onboarding + walkthrough + trusted-contacts + welcome view · English, i18n-ready.
- **Tech:** React Native + Expo · Supabase · Edge-Function dispatch · Expo Push · free-GPS + external maps · Next.js admin.

**Explicitly DEFERRED (see phase tags throughout):**

- **[P2]:** Premium Choice (seeker picks helper) · masked calling · SOS Layer 3 (community responders) · E2E inbox with client-side reporting · retroactive help-logging · voice-input category analysis · the anti-collusion layer · real map embed · real-world Moneta redemption (*gated behind the anti-collusion layer + regulatory review*).
- **[LATER]:** formal background checks · behavioural trust gates · direct DigiLocker integration · cash rewards & City Saviour roles (human-judged) · Hindi/regional languages · disappearing messages · arbitrary group chats · comments.
- **[GATE — blocks public launch]:** legal entity, liability/ToS, DPDP, POC SO, IT Rules, insurance, Moneta-redemption regulation (§2.4 G1–G7); App Store review (§11.10); DLT SMS registration (§11.7).

---

**OPEN ITEMS TRACKER (rolling — tuning + parked items)**

| ID | Item                                                                                                                                                                                          | Bucket         | Status                |
|----|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------|-----------------------|
| A1 | <del>Urgency self-declared vs enforced.</del><br><b>RESOLVED</b> (10.5):<br>pre-filled from category<br>typicalUrgency default<br>+ light user nudge; SOS excluded (separate guarded button). | 0.2 / 3 / 10.5 | ✅ Closed              |
| A2 | Moderation <b>operations</b><br>— who reviews                                                                                                                                                 | 1.4            | Open — dedicated pass |

| ID | Item                                                                                                                                                                                    | Bucket         | Status                                   |
|----|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------|------------------------------------------|
|    | reports, escalation, response-time targets, 11pm reachability, code of conduct, dispute resolution.                                                                                     |                |                                          |
| A3 | New-helper / cold-start-of-trust.<br><i>Partially addressed:</i> badge doesn't affect dispatch order (3.4), so new helpers get pinged fairly; seeker-skips-new-moon-at-confirm remains. | 1.6 / 3.4      | Open (narrowed) — density bucket         |
| A4 | Choose KYC vendor (needs research).                                                                                                                                                     | 2.2            | Open                                     |
| A5 | <del>What does Moneta redeem for?</del><br><b>RESOLVED</b> (7.3): recognition/access in v1, real-world value LATER behind the anti-abuse gate (G7).                                     | 7.3            | ✅ Closed                                 |
| A6 | <b>Funding the organization</b> — grants, CSR, donations, impact investment, Premium Choice. The biggest open question.                                                                 | Org / §12.3    | 🛑 Parked — dedicated session (conscious) |
| A7 | <del>Cold-start density.</del><br><b>RESOLVED</b> (§12.1): launch via <b>closed communities</b> (colleges/companies/orgs) — solves density, cold-start, and baseline trust together.    | Growth / §12.1 | ✅ Closed                                 |
| A8 | Accessibility & the low-tech user; multilingual India.                                                                                                                                  | UX / 10        | Open (narrowed)                          |



| ID  | Item                                                                                                                                                                                 | Bucket              | Status          |
|-----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------|-----------------|
|     | <i>Partially addressed:</i><br>70-year-old test drives UI (10.1), accessible-by-default, big tiles, voice-input P2; <b>English first, i18n-ready for Hindi/regional next</b> (10.2). |                     |                 |
| A9  | <del>Notification strategy</del> — <b>RESOLVED</b> in 3.3 / 3.8 / 3.9 (waves, soft daily cap + SOS exempt, digest, seeker-notify moments).                                           | 3                   | ✓ Closed        |
| A10 | Data retention (location history, chat), law-enforcement requests. <i>Partially addressed:</i> chat retention decided (4.3); location history + LE requests remain.                  | Privacy             | Open (narrowed) |
| A11 | <del>Success metrics</del> . <b>RESOLVED</b> (§12.2): north-star = completed <b>unique helps</b> ; guardrails = safety incidents ( $\approx 0$ ) + helper retention.                 | Measurement / §12.2 | ✓ Closed        |
| A12 | Wave sizes + intervals per category/urgency — tune with real usage.                                                                                                                  | 3.3                 | Open — tuning   |
| A13 | Now-request default lifetime windows per tier — tune.                                                                                                                                | 3.7                 | Open — tuning   |
| A14 | Notification daily-cap number + digest cadence — tune.                                                                                                                               | 3.8                 | Open — tuning   |

| ID  | Item                                                                                                                                      | Bucket    | Status                    |
|-----|-------------------------------------------------------------------------------------------------------------------------------------------|-----------|---------------------------|
| A15 | Chat data retention window (30–90d?) + media file-size caps — set with counsel (DPDP) and tune.                                           | 4.3 / 4.4 | Open — tuning + [GATE G3] |
| A16 | <b>Write the code of conduct</b> — the banned-behaviour list the reporting loop enforces. Needed before launch.                           | 4.5 / 1.4 | Open                      |
| A17 | Directed-request fallback window — how long before it reverts to open dispatch. Tune.                                                     | 5.5       | Open — tuning             |
| A18 | <b>Measure leakage</b> (informal help between connections) post-launch — the data that decides whether retroactive help-logging is built. | 5.11      | Open — needs live data    |
| A19 | Inbox media expiry window (~6–12 months) — set with storage cost + DPDP in mind.                                                          | 6.7       | Open — tuning             |
| A20 | Active-request chat rate-limit thresholds — tune so normal users never notice them.                                                       | 6.12      | Open — tuning             |
| A21 | Goodness Meter curve constant <i>k</i> — tune so the needle feels responsive for a normal helper.                                         | 7.7       | Open — tuning             |
| A22 | Leaderboard is auto-opt-in — revisit if                                                                                                   | 7.9       | Open — monitor            |

| ID  | Item                                                                                     | Bucket | Status          |
|-----|------------------------------------------------------------------------------------------|--------|-----------------|
|     | quiet helpers want to stay off it.                                                       |        |                 |
| A23 | Final font sign-off (Nunito Sans workhorse; Space Grotesk held as headings-only option). | 10.2   | Open — decision |
| A24 | Choose SMS provider (needs live research on current pricing/deliverability).             | 11.7   | Open            |

---

#### DOCUMENT STATUS — COMPLETE

**Sapiens is fully designed.** Eleven buckets, every mechanic decided, every major tension resolved, all with rationale recorded and phase-tagged.

#### What remains is not product design:

1. **Deliberately deferred features** — P2/Later, each with a v1 architectural hook (§13).
2. **Tuning values** — set with real usage, not guessed (Open Items Tracker: A12–A24).
3. **Pre-launch gates** — legal (G1–G7), App Store review, DLT registration; resolved before public launch.
4. **Parked strategic question** — funding (§12.3), its own session.

**Next deliverable: the MVP Build Specification** — this PRD distilled to **P1 only**, written for Claude Code, with the architecture made concrete (Supabase schema + RLS, the dispatch-engine spec, a staged build order). The PRD is the *what and why*; the build spec is the *what to build first and how*. Kept as **two separate documents** — the PRD is the timeless design record; the build spec is the lean, disposable instruction set for this build.